


CPU

1690 pJ/flop

Optimized for Latency
Caches




Westmere
32 nm

GPU

140 pJ/flop

Optimized for Throughput
Explicit Management
of On-chip Memory



Kepler
28 nm